

Tidy Inference in R

Data Science Programming

Alison Hill

2018-03-09

Install Packages

```
install.packages("tidyverse")
install.packages("broom")
# today I'll use the development version
install.packages("remotes")
remotes::install_github("andrewpbray/infer", ref = "develop")
install.packages("devtools")
devtools::install_github("sfirke/janitor")
```

Load Packages

```
library(tidyverse)  
library(infer)  
library(janitor)  
library(broom)
```



Research Question

If you see someone else yawn, are you more likely to yawn?

In an episode of the show **Mythbusters**, they tested the myth that yawning is contagious.

Participants and Procedure

Participants and Procedure

- $N = 50$ adults who thought they were being considered for an appearance on the show.

Participants and Procedure

- $N = 50$ adults who thought they were being considered for an appearance on the show.
- Each participant was interviewed individually by a show recruiter ("confederate") who either
 - yawned, $n = 34$
 - or not, $n = 16$.

Participants and Procedure

- $N = 50$ adults who thought they were being considered for an appearance on the show.
- Each participant was interviewed individually by a show recruiter ("confederate") who either
 - yawned, $n = 34$
 - or not, $n = 16$.
- Participants then sat by themselves in a large van and were asked to wait.

Participants and Procedure

- $N = 50$ adults who thought they were being considered for an appearance on the show.
- Each participant was interviewed individually by a show recruiter ("confederate") who either
 - yawned, $n = 34$
 - or not, $n = 16$.
- Participants then sat by themselves in a large van and were asked to wait.
- While in the van for some set amount of time (unknown), the Mythbusters watched to see if the unaware participants yawned.

Data

Two group design:

- $n = 34$ saw the confederate yawn (*seed*)
- $n = 16$ did not see the confederate yawn (*control*)

Data

Two group design:

- $n = 34$ saw the confederate yawn (*seed*)
- $n = 16$ did not see the confederate yawn (*control*)

```
group <- c(rep("control", 12), rep("seed", 24),  
          rep("control", 4), rep("seed", 10))  
yawn <- c(rep(0, 36), rep(1, 14))  
  
yawn_myth <- data_frame(subj = seq(1, 50), group, yawn) %>%  
  mutate(yawn = as.factor(yawn))  
  
glimpse(yawn_myth)
```

Observations: 50

Variables: 3

```
$ subj <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 1 ...  
$ group <chr> "control", "control", "control", "control", "control", " ...  
$ yawn <fct> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
```

Show entries

Search:

	subj	group	yawn
1	1	control	0
2	2	control	0
3	3	control	0
4	4	control	0
5	5	control	0
6	6	control	0
7	7	control	0
8	8	control	0

Showing 1 to 8 of 50 entries

Previous

1

2

3

4

5

6

7

Next

Results

```
yawn_myth %>%  
  tabyl(group, yawn) %>%  
  adorn_percentages() %>%  
  adorn_pct_formatting() %>%  
  adorn_ns()
```

group	0	1
control	75.0% (12)	25.0% (4)
seed	70.6% (24)	29.4% (10)

Conclusion

Conclusion

Finding: CONFIRMED¹

[1] <http://www.discovery.com/tv-shows/mythbusters/mythbusters-database/yawning-contagious/>

Really?

"Though that's not an enormous increase, since they tested 50 people in the field, the gap was still wide enough for the MythBusters to confirm that yawning is indeed contagious."¹

[1] <http://www.discovery.com/tv-shows/mythbusters/mythbusters-database/yawning-contagious/>

State the hypotheses

State the hypotheses

H_0 :

There is no difference between the seed and control groups in the proportion of people who yawned.

State the hypotheses

H_0 :

There is no difference between the seed and control groups in the proportion of people who yawned.

H_1 (directional):

More people (relatively) yawned in the seed group than in the control group.

Test the hypothesis

Which type of hypothesis test would you conduct here?

- Independent samples t-test
- Two proportion test
- Chi-square test of independence
- Analysis of Variance
- I don't know!

Test the hypothesis

Which type of hypothesis test would you conduct here?

- Independent samples t-test
- Two proportion test
- Chi-square test of independence
- Analysis of Variance
- I don't know!

Test the hypothesis

Which type of hypothesis test would you conduct here?

- Independent samples t-test
- Two proportion test
- Chi-square test of independence
- Analysis of Variance
- I don't know!

Answer:

- Two proportion test

Two proportion test

Two proportion test

$$H_0 : p_{seed} = p_{control}$$

Two proportion test

$$H_0 : p_{seed} = p_{control}$$

$$H_1 : p_{seed} > p_{control}$$

The observed difference

```
yawn_myth %>%  
  group_by(group) %>%  
  summarize(prop = mean(yawn == 1))
```

```
# A tibble: 2 x 2  
  group    prop  
  <chr>   <dbl>  
1 control 0.250  
2 seed    0.294
```

The observed difference

```
yawn_myth %>%  
  group_by(group) %>%  
  summarize(prop = mean(yawn == 1))
```

```
# A tibble: 2 x 2  
  group    prop  
  <chr>   <dbl>  
1 control 0.250  
2 seed    0.294
```

```
(obs_diff ← yawn_myth %>%  
  group_by(group) %>%  
  summarize(prop = mean(yawn == 1)) %>%  
  summarize(diff(prop)) %>%  
  pull())
```

```
[1] 0.04411765
```

Is this difference *meaningful*?

Is this difference *meaningful*?

Different question:

Is this difference *meaningful*?

Different question:

Is this difference *significant*?

Modeling the null hypothesis

If...

$$H_0 : p_{seed} = p_{control}$$

is true, then whether or not the participant saw someone else yawn does not matter: there is no association between exposure and yawning.



Original universe

```
# A tibble: 12 x 3
  subj group   yawn
<int> <chr>   <fct>
1     1   1 control 0
2     2   2 control 0
3     3   3 control 0
4     4   4 control 0
5     5   5 control 0
6     6   6 control 0
7    15  15 seed    0
8    16  16 seed    0
9    17  17 seed    0
10   18  18 seed    0
11   19  19 seed    0
12   20  20 seed    0
```

group	0	1	Total
control	12	4	16
seed	24	10	34
Total	36	14	50

Original universe

```
# A tibble: 12 x 3
  subj group yawn
<int> <chr> <fct>
1     1   control 0
2     2   control 0
3     3   control 0
4     4   control 0
5     5   control 0
6     6   control 0
7    15   seed    0
8    16   seed    0
9    17   seed    0
10   18   seed    0
11   19   seed    0
12   20   seed    0
```

group	0	1	Total
control	12	4	16
seed	24	10	34
Total	36	14	50

Parallel universe

```
# A tibble: 12 x 3
  subj group alt_yawn
<int> <fct> <fct>
1     1   control 0
2     2   control 0
3     3   control 1
4     4   control 0
5     5   control 0
6     6   control 0
7    15   seed    0
8    16   seed    1
9    17   seed    0
10   18   seed    0
11   19   seed    0
12   20   seed    1
```

group	0	1	Total
control	12	4	16
seed	24	10	34
Total	36	14	50

1000 parallel universes


```
# A tibble: 1,000 x 2
  replicate    stat
  <int>      <dbl>
1         1 -0.140
2         2 -0.0478
3         3 -0.0478
4         4 -0.0478
5         5 -0.140
6         6  0.228
7         7 -0.140
8         8  0.136
9         9 -0.232
10        10  0.0441
# ... with 990 more rows
```

1000 parallel universes

```
# A tibble: 1,000 x 2
  replicate    stat
  <int>      <dbl>
1         1 -0.140
2         2 -0.0478
3         3 -0.0478
4         4 -0.0478
5         5 -0.140
6         6  0.228
7         7 -0.140
8         8  0.136
9         9 -0.232
10        10  0.0441
# ... with 990 more rows
```

```
# A tibble: 11 x 2
  replicate    stat
  <int>      <dbl>
1       990 -0.232
2       991 -0.0478
3       992 -0.232
4       993  0.136
5       994 -0.140
6       995  0.136
7       996 -0.232
8       997  0.0441
9       998  0.0441
10      999 -0.232
11     1000 -0.140
```

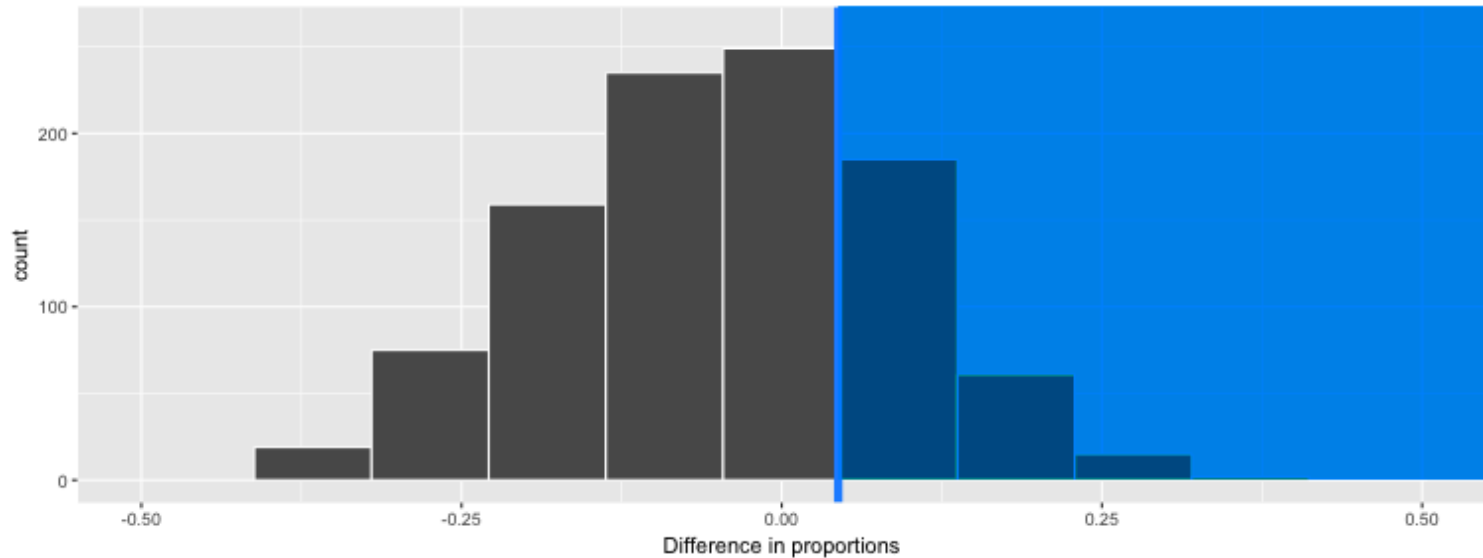
The parallel universe distribution



The distribution of 1000 differences in proportions, if the null hypothesis were *true* and yawning was not contagious.

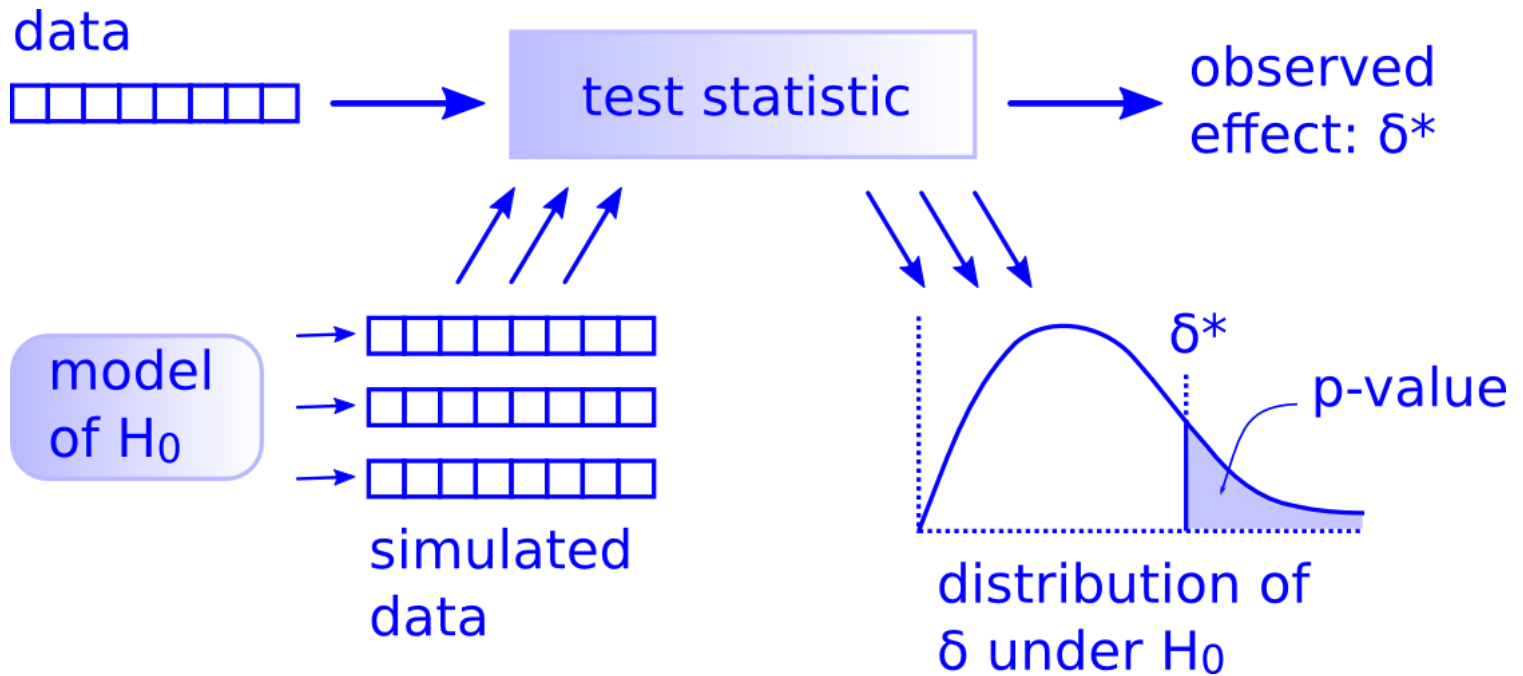
In how many of our "parallel universes" is the difference as big or bigger than the one we observed (0.0441176)?

Calculating the p-value



That proportion is the p-value!

```
# A tibble: 1 x 3
  n_as_big n_total p_value
  <int>    <int>    <dbl>
1     512    1000    0.512
```



The tidy way
Use the infer package

Hypothesis test




Hypothesis test




`specify()`

<https://github.com/ismayc/talks/tree/master/data-day-texas-infer>

Hypothesis test



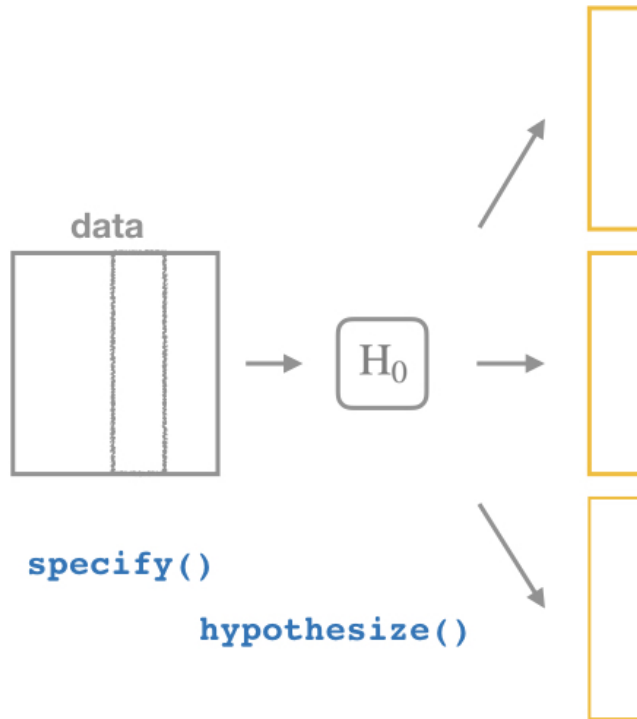
Hypothesis test



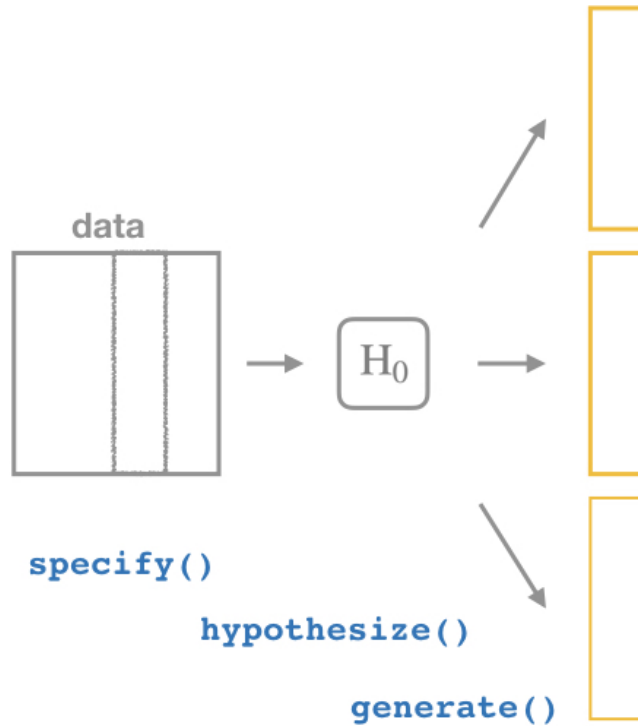
`specify()`

`hypothesize()`

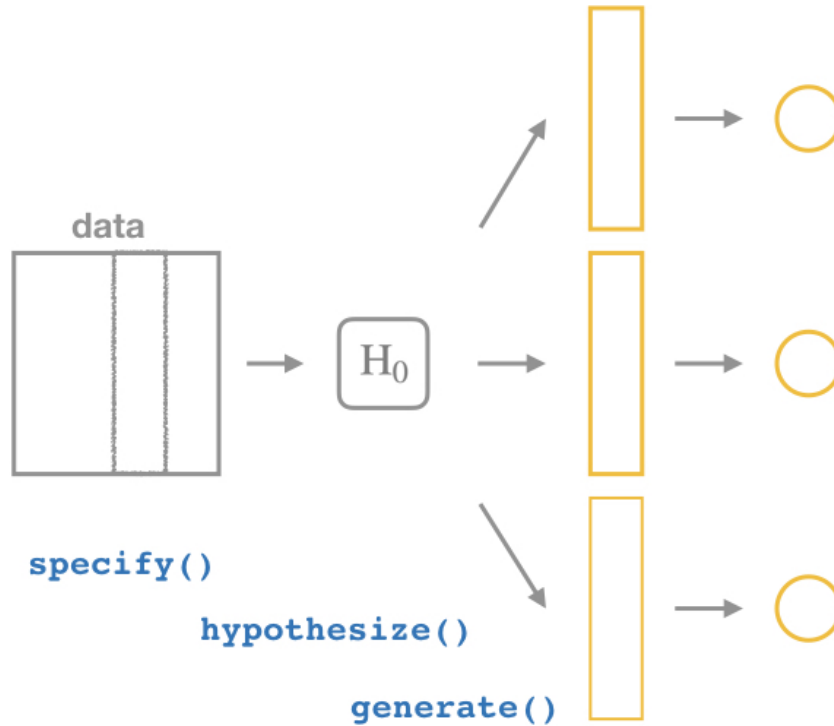
Hypothesis test



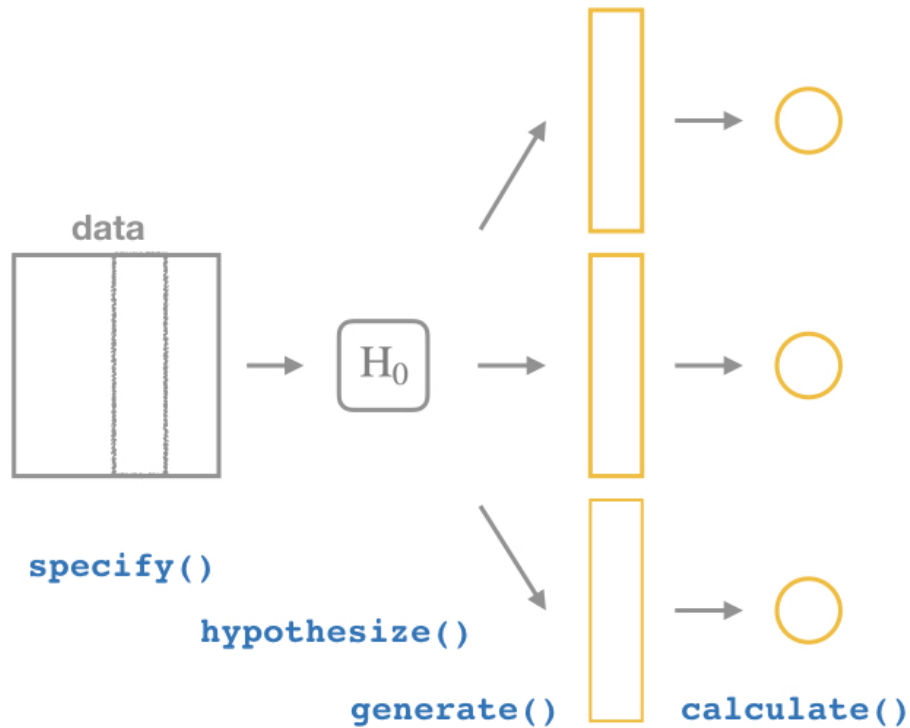
Hypothesis test




Hypothesis test




Hypothesis test




Hypothesis test



Hypothesis test



Confidence Interval



infer

5 functions:

- `specify()`
- `hypothesize()`
- `generate()`
- `calculate()`
- `visualize()`


infer

- **specify** the response and explanatory variables ($y \sim x$)
- **hypothesize** what the null H_0 is (here, **independence** of y and x)
- **generate** new samples from parallel universes:
 - Resample from our original data **without replacement**, each time shuffling the **group** (`type = "permute"`)
 - Do this **a ton of times** (`reps = 1000`)
- **calculate** the *new* statistic (`stat = "diff in props"`) for each **rep**

```
set.seed(8)
null_distn <- yawn_myth %>%
  specify(formula = yawn ~ group, success = "1") %>%
  hypothesize(null = "independence") %>%
  generate(reps = 1000, type = "permute") %>%
  calculate(stat = "diff in props", order = c("seed", "control"))
```

Visualize the null distribution

- `visualize` the distribution of the `stat` (here, `diff in props`)



```
visualize(null_distn, bins = 10) +  
  geom_vline(xintercept = obs_diff,  
             color = "dodgerblue",  
             size = 2)
```

Classical inference

Rely on theory to tell us what the null distribution looks like.

```
yawn_myth %>%  
  specify(yawn ~ group, success = "1") %>%  
  hypothesize(null = "independence") %>%  
  # generate() is not needed since we are not doing randomization  
  # calculate(stat = "z") is implied based on variable types  
  visualize(method = "theoretical") +  
  geom_vline(xintercept = obs_stat, color = "orchid", size = 2)
```


Classical vs resampling

Changed the `stat` to calculate to `z` now (before we did `diff in props`).

```
yawn_myth %>%  
  specify(yawn ~ group, success = "1") %>%  
  hypothesize(null = "independence") %>%  
  generate(reps = 1000, type = "permute") %>%  
  calculate(stat = "z", order = c("seed", "control")) %>%  
  visualize(method = "both", bins = 10) +  
  geom_vline(xintercept = obs_stat, color = "orchid", size = 2)
```

Do the test in R

```
yawn_table <- table(group, yawn)
yz <- prop.test(x = yawn_table, n = nrow(yawn_myth),
               alternative = "greater", correct = FALSE)
yz
```

2-sample test for equality of proportions without continuity correction

```
data:  yawn_table
X-squared = 0.10504, df = 1, p-value = 0.3729
alternative hypothesis: greater
95 percent confidence interval:
 -0.1754872  1.0000000
sample estimates:
   prop 1    prop 2 
0.7500000 0.7058824
```

Do the test in R

```
yawn_table <- table(group, yawn)
yz <- prop.test(x = yawn_table, n = nrow(yawn_myth),
               alternative = "greater", correct = FALSE)
yz
```

2-sample test for equality of proportions without continuity correction

```
data:  yawn_table
X-squared = 0.10504, df = 1, p-value = 0.3729
alternative hypothesis: greater
95 percent confidence interval:
 -0.1754872  1.0000000
sample estimates:
   prop 1    prop 2 
0.7500000 0.7058824
```

What does this test assume?

Pull out the z statistic

The observed `z` value is 0.105042.

```
obs_stat <- yz %>%  
  broom::tidy(yz) %>%  
  pull(statistic)
```

Now you can use `obs_stat` in `geom_vline(xintercept = obs_stat)` when added to `visualize!`

In class exercise

- Read in the `mazes` data (<http://bit.ly/mazes-gist>)

```
library(readr)
mazes <- read_csv("http://bit.ly/mazes-gist") %>%
  clean_names() %>% #janitor package
  filter(dx %in% c("ASD", "TD"))
```

- Use `dplyr::filter` to include only two groups (`DX` if you didn't `clean_names()`; `dx` if you did!).
- Use `infer` to compare a numerical variable between the two groups using:
 - A permutation test and
 - A classical theoretical test.

See: <https://infer-dev.netlify.com>

About the data: [Quantitative analysis of disfluency in children with autism spectrum disorder or language impairment](#)

Classical t-test in R

- Independent samples (`paired = FALSE`)
- Assume equal variances (`var.equal = TRUE`)
- `alternative` is relative to the groups *alphabetically*: so here $H_1 = \mu_{asd} < \mu_{td}$

```
myt <- t.test(viq ~ dx,  
              data = mazes,  
              var.equal = TRUE,  
              paired = FALSE,  
              alternative = "less")  
  
myt
```

Two Sample t-test

```
data: viq by dx  
t = -11.842, df = 308, p-value < 2.2e-16  
alternative hypothesis: true difference in means is less than 0  
95 percent confidence interval:  
 -Inf -18.63839  
sample estimates:  
mean in group ASD  mean in group TD  
    95.28962      116.94488
```

Save the t statistic

```
obs_t <- myt %>%  
  tidy() %>% # from broom  
  pull(statistic) # from dplyr  
obs_t
```

```
[1] -11.84247
```

Now you can use `obs_t` in `geom_vline(xintercept = obs_t)` when added to `visualize!`

Thanks!

Slides created via the R package [xaringan](#).